

Spack Work Items for HEP

Chris Green greenc@fnal.gov Lynn Garren garren@fnal.gov
Patrick Gartung gartung@fnal.gov

version 0.4 (2019-05-16)

Contents

1	Introduction	3
2	High-Impact Core Enhancements	3
2.1	Support for external Spack commands	3
2.1.1	Basic Support for external Spack commands	3
2.1.2	Enhancement of Spack subcommand facilities	3
2.1.3	Support for flake8 checks of external Spack commands . .	3
2.1.4	Support for external commands and support code in external package repos	3
	Status	3
2.2	The ability to layer Spack installations in a hierarchy	4
	Status	4
2.3	A persistent relationship between a built package and the options used to build it	4
	Status	4
2.4	Support for compiler-agnostic packages	5
	Status	5
2.5	No-source packages	5
	Status	5
2.6	Base build location on stage rather than expanded source location	6
	Status	6
2.7	<code>self.stage.source_path</code> should be usable before it exists. . . .	6
	Status	6
2.8	Masking of installed packages and dependents	6
	Status	6
3	Other Spack Core Enhancements	7
3.1	The ability to inherit from package-repo-specific package base classes	7
	Status	7

3.2	New environment manipulation functions for system path de- prioritization and duplicate removal	7
	Status	7
3.3	Configurable equivalency between operating systems	7
	Status	8
3.4	Updates to the version comparison system	8
3.4.1	Support correct ordering of alpha, beta, pre, rc and patch versions	8
3.4.2	Support for half-open and closed intervals for version ranges	8
3.4.3	Correct ordering between versions X and X.Y	8
3.4.4	<code>version()</code> should recognize filename extensions from <code>url_from_version()</code> including <code>tbz2</code>	9
	Status	9
3.5	Compact variant propagation via <code>depends_on()</code>	9
	Status	10
3.6	Improvements to the variants system	10
3.6.1	<code>variant()</code> grows a <code>when</code> clause	10
3.6.2	New function <code>variant.depends_on()</code>	10
	Status	10
3.7	Install only specified packages from YAML specifications.	10
	Status	11
3.8	Writing of the YAML spec file at the start of the build process rather than solely at install time.	11
	Status	11
4	Recipe Enhancements	11
4.1	Client-only mysql-ish packages	11
	Status	11
4.2	Client-only postgresql-build	11
	Status	11
4.3	Updated root recipe	12
	Status	12
4.4	Propagate C++ attributes through dependencies of C++-using packages	12
	Status	12
5	Version History	12
	Version 0.1 (2019-02-25)	12
	Version 0.2 (2019-03-19)	12
	Version 0.3 (2019-04-30)	13
	Version 0.4 (2019-05-16)	13
	Version 0.5 (2019-06-11)	13

1 Introduction

While working on the Spack-based build, development, and release system for HEP, we have come across many instances big and small where Spack core or its recipes require enhancements to allow us to meet our project requirements. In some cases these are showstoppers, such as support for multiple releases and layered (central vs local) installations; in others, they improve an aspect of Spack operation without being essential, such as providing repo-specific package base classes.

This document is an attempt to be a descriptive list of these work items (the most impactful and necessary core items are listed first), along with (where appropriate) ideas for implementation, statement of importance, and current status.

2 High-Impact Core Enhancements

2.1 Support for external Spack commands

2.1.1 ~~Basic Support for external Spack commands~~

In order to support the efficient operation of Spackdev, allowing the use of Spack functions and structures internally, `spack dev` should be a spack command. For purposes of ownership and ease of maintenance, this would ideally be an external command, analogous to (*e.g.* `git flow` as an external `git` command).

2.1.2 Enhancement of Spack subcommand facilities

Spack's existing handling of subcommands should be generalized to work for external commands and at any level.

2.1.3 Support for flake8 checks of external Spack commands

A simple way to check external Spack command code for rule-correctness is strongly desired.

2.1.4 Support for external commands and support code in external package repos

The ability to find external commands in external package repos. This is related to a [Spack Core Enhancement](#) described elsewhere herein.

Status

[PR 8612](#) from Massimiliano Culpo implemented basic support for external Spack commands. It is now integral to the operation of SpackDev and was merged into the FNAL fork on 2018-09-19. The PR was merged upstream on 2019-03-28

with some changes from the original PR submission as merged into the FNAL fork, which was subsequently reconciled with upstream.

A git branch consisting solely of [our proposed enhancement of the subcommand facilities](#) was submitted as [PR 11145](#). Comments were received and incorporated. Subsequently, a competing [PR 11209](#) was submitted by the original author of [PR 8612](#) purporting to address one of the issues addressed by [PR 11145](#). This PR was incorporated into [PR 1145](#) and is awaiting further comments.

While [PR 8612](#) supports testing of external Spack commands, it appears not to support intuitive coding rule enforcement for same. An enhancement request should be discussed with the author of [PR 8612](#).

2.2 ~~The ability to layer Spack installations in a hierarchy~~

In order to be able to combine locally-built and centrally-installed packages, it becomes necessary to be able to refer to packages in a different Spack installation from the current one.

Status

Originally conceived by Jim Amundson (as “Spack Chains”) and entered as [PR 5014](#), and subsequently resubmitted against a more recent version of Spack development by Patrick Gartung as [PR 8545](#), this feature has been taken up by Peter Scheibel—a Spack principal— as [PR 8772](#). This PR was merged on 2019-03-27 and a subsequent [PR 11152](#) adding documentation was merged 2019-04-10.

2.3 A persistent relationship between a built package and the options used to build it

In order to support the use of a given Spack installation to support multiple installed software distributions over a significant period of time, it is necessary to be able to identify an installed or cached binary package such that it remains possible to use it even after updates to the recipe(s) have caused a change to the already-built package’s calculated hash.

It is the nature of HEP software installations to have a relatively static central installation of software packages and distribution releases against which higher-level packages are developed by users. It must be possible to install centrally newer distribution releases while not compromising the ability of users to build and develop their own software against older ones without unnecessary rebuilds of lower level packages due to hash changes resulting from updated recipes.

Status

The recent [Spack Environments](#) functionality may assist with this (as may [Spack Chains](#), described elsewhere in this document), but this remains a thorny problem

that will require discussion and exploration before a solution is implemented. The subject has been discussed with Spack principals such as Todd Gamblin and Peter Scheibel and we have been assured that this feature is compatible with the philosophy and direction of Spack, that it has other desirous parties than us, and that it would be reasonable to expect something to look at that should do the job by the end of September 2019.

2.4 Support for compiler-agnostic packages

The HEP community has several applications, distributions of which must be accompanied by substantial data. Examples are particle interaction descriptions and cross-section table. Examples outside our field might include star catalogs and other astrophysics data. Information of this nature can be quite voluminous ($\geq 5\sim\text{GiB}$). We would wish to avoid unnecessary duplication for multiple compilers and versions by having just one compiler-agnostic package containing the data.

In addition, pure C and pure interpreted-language packages such as `fftw` or `pyyaml`—or those which combine same—are similarly independent of compiler and version thereof (or perhaps require only a generic compiler runtime such as `libgcc_s.o`). A way to avoid unnecessary duplication of these packages also would be desirable, although the win here would be less in space and more in CPU time for builds.

Status

The subject has been discussed with Spack principals such as Todd Gamblin and Peter Scheibel, including at the telecon 2019-04-11, but no formal enhancement request has been made. Something providing this functionality is expected to be part of the upcoming concretization improvements, but it should be noted that merely allowing the compiler as a “normal” dependency is insufficient to provide the needed flexibility.

2.5 No-source packages

In our existing system, we make significant use of no-source, or “umbrella” packages, whose only function is to facilitate the build or setup of a particular collection of packages.

Status

This feature has been discussed with Spack principals. An available workaround is the use of a webpage URL as a “source.”

~~2.6 Base build location on stage rather than expanded source location~~

The definition of the CMake build area should be changed from being relative to the staged source directory to being relative to the stage directory.

We have use cases where a recipe needs to have information about the build directory in the environment set by `setup_environment()` (*e.g.* if tests need to have a directory in `PATH`). Specifically, staging happens *after* the call to `setup_environment()`, and if the stage area does not exist, then `spec.stage.source_path` returns, `None`.

Status

[PR 8431](#) was submitted 2018-06-08 and was moribund for several months. This PR was eventually accepted and merged upstream on 2019-03-26.

2.7 `self.stage.source_path` should be usable before it exists.

There are circumstances where it is desirable to *e.g.* set an environment variable in `setup_environment()` based on `self.stage.source_path`. Currently this is not possible as its existence is enforced and `setup_environment()` is executed prior to any staging of the source.

Status

Currently workarounds are used when necessary: adding to the default environment for the `make()` object, for example. PRs [11528](#) (merged upstream) and [11662](#) (pending) purport to address this issue, but have not yet been examined.

2.8 Masking of installed packages and dependents

A user may wish to develop packages that may already be installed (centrally, perhaps). We require a mechanism to hide these already-installed packages and their dependents from the concretization system

Status

This feature is not currently required for the Minimum Viable Product (MVP), but will be essential in the finished system. It is possible that environments or views may satisfy this requirement, although that remains to be investigated.

3 Other Spack Core Enhancements

3.1 The ability to inherit from package-repo-specific package base classes

Fermilab has a large number of experiment-related packages using a common CMake-based system called `cetmodules`. Recipes for these packages share a significant amount of boilerplate which could profitably be part of a package base class situated in the same package repo.

Status

This feature has been discussed informally with Spack principals.

3.2 New environment manipulation functions for system path deprioritization and duplicate removal

In order to prevent contamination of `PATH`-like variables with system paths due to `packages.yaml` entries, several new functions provide facilities for `PATH`-like variable manipulation.

Status

[PR 8476](#) was submitted 2018-06-14 and was moribund for a long time. It has been updated against upstream and discussed with Spack principals and further information has been added to the PR. After changes were made to accommodate upstream developments, it was merged on 2019-05-06.

[PR 11434](#) was submitted 2019-05-13 to allow the production of source-able text and pickled environment files via the command line (`spack build-env --pickle <file>`, `spack build-env --dump <file>`). Initial comments were addressed and the PR is awaiting final approval and merge.

3.3 Configurable equivalency between operating systems

Due to a change in how Scientific Linux installations identify themselves between SL7.5 and SL7.6, there needs to be a configurable equivalency between compatible operating systems *e.g.* in `config.yaml`:

```
equivalent-os:
  rhel6:
    - scientific6
    - centos6
  rhel7:
    - scientific7
    - centos7
```

Status

We have a hard-wired equivalency between `rhel7` and `scientific7` in our local code.

A concept for the solution exists (see above). Analysis is needed to determine feasibility and consequences for (*e.g.*) hash values, buildcache packages, *etc.*

This issue has been discussed with Spack principals: OS string conversion, and updating the metadata system to handle equivalency. This issue also has relevance to CPU architecture (Haswell, *etc.*).

3.4 Updates to the version comparison system

3.4.1 Support correct ordering of alpha, beta, pre, rc and patch versions

While it is not usual to need to refer to non-mainstream versions in a recipe, it is certainly not beyond imagining. Several packages have had long release candidate cycles, for instance, and a build of one might be necessary to address an observed issue prior to the final release becoming public. When non-mainstream versions are used in ranges, the version comparison needs to be aware of the relative ordering—release candidates following “pre” versions, but preceding their corresponding released versions, for example.

3.4.2 Support for half-open and closed intervals for version ranges

e.g. describing something which is valid for $7.0 \leq \text{version} < 8.0$. For purely numeric versions, the usual workaround of `when=@7.0:7.999` suffices, but for something which is valid also for (*e.g.* release candidates), it becomes more desirable to be able to specify “everything before ...” Example nomenclature:

```
# No support before ...
conflicts('cxxstd=17', when='@:<1.27.0')
# Blacklist range.
variant('bad_idea', default=True, when='@0.5.0:<1.0.0')
# Last known good.
conflicts('went_bad_after', when='@1.5.7>:')
# Ignore 2.0.0-alpha, pre, etc.
conflicts('Shaky prerelease', when='@1.8.12>:<@2.0.0')
```

3.4.3 Correct ordering between versions X and X.Y

There appears to be a problem ordering versions with different levels of specification. For instance, in `intel-tbb`, the version version of any year’s releases is the unadorned year (2019, say), with 2019U1 being the next. These have to be encoded currently as 2019.0 and 2019.1, because 2019 does not appear to compare less-than 2019.1. Following discussion with Peter Scheibel, it appears that this is intentional behavior in the numeric case, and additionally applies to

non-numeric versions where one is a front-anchored substring of another (*e.g.* MVP1 vs MVP1a).

3.4.4 `version()` should recognize filename extensions from `url_from_version()` including `tbz2`

Several of our recipes are of the form where the following `url_from_version()` is appropriate:

```
def url_for_version(self, version):
    url = 'http://cdcvs.fnal.gov/cgi-bin/git_archive.cgi/' \
          'cvs/projects/{0}.v{1}.tbz2'
    return url.format(self.name, version.underscored)
```

Unfortunately, for every call to `version()`, it is necessary to specify the extension explicitly:

```
version('2.3.0',
        sha256='4b6a29443b631957ca2a7712b5c57762'
        '0c6543e542ee9c77d246cef1e10f7324',
        extension='tbz2')
```

The system should be trained to understand `tbz2` as a valid extension when returned as a URL by `url_for_version()`.

Status

These core enhancements are at the conceptual stage only. No implementation has been attempted, but it is likely that they are perfectly feasibly implementable by a contributor rather than a core expert. This could be entered instead as an enhancement request, but the likelihood of a core developer taking it on is somewhat low. The truncated-version-ordering issue must be addressed with care as there are use cases for both behaviors, and they should be distinguishable and selectable as appropriate.

3.5 Compact variant propagation via `depends_on()`

There is a need for propagation of variants to dependencies generally, rather than specifically. For instance, in order to pass the value of the `cxxstd` variant down a dependency tree, currently one must do:

```
depends_on('boost cxxstd=default', when='cxxstd=default')
depends_on('boost cxxstd=98', when='cxxstd=98')
...
```

One could imagine instead a more compact syntax, something like:

```
depends_on('boost', propagate_variants=['cxxstd', 'python'])
```

If `cxstd` is not specified, then nothing will be propagated (if the variant exists in the dependency it will be defaulted). Otherwise, the value of `cxstd` in the current recipe will be used.

As a further extension, using `recursive_variants` instead of `propagate_variants` would cause the specified variants to be propagated to lower-level dependencies *even if intermediate dependencies did not have the variant(s) in question defined*.

Status

These enhancements are at the conceptual stage only.

3.6 Improvements to the variants system

Several packages have extensive feature options, which evolve over time. Features come and go, and depend upon each other. An expanded variant specification syntax is desired to be able to specify these variants and the relationships between them. The archetype for this use case would be the `boost` recipe. Some ideas:

3.6.1 `variant()` grows a `when` clause

The `when` clause shall specify when the variant is defined. It shall be valid therein to refer to versions and to any other variants that have already been defined.

3.6.2 New function `variant.depends_on()`

This function, having a `when` clause as defined for `variant()`, above, shall specify upon what other variant(s) `self` shall depend when the `when` clause is valid. It shall return the variant object allowing for chaining, *e.g.*

```
variant('context', when='@1.61.00:').\
  depends_on('thread', 'chrono', 'system', 'date_time')
variant('fiber', when='@1.62.00:').depends_on('context').\
  depends_on('thread', 'chrono', 'system', 'date_time',
            when='@:1.68.99')
variant('signals', when='@1.29.00:1.68.99')
```

Status

These enhancements are at the conceptual stage only.

3.7 Install only specified packages from YAML specifications.

Spack currently provides the ability to install packages specified by a YAML spec file in addition to those specified on the command line. Spack should provide the ability to install the *intersection* of these two classes of package specification

rather than their union. This improves the ability to easily (re-)install subsets of a large and complex spec tree.

Status

These enhancements are at the conceptual stage only.

3.8 Writing of the YAML spec file at the start of the build process rather than solely at install time.

To improve the development cycle time for packages that are part of complex dependency hierarchies, the YAML spec tree should be written in the same directory as `spack-build.{env,out}` to allow error-free repeats of install attempts.

Status

These enhancements are at the conceptual stage only.

4 Recipe Enhancements

4.1 ~~Client-only mysql-ish packages~~

The existing `mariadb` package is GPL and therefore not usable by HEP community in a linking context. An LGPL client-only package provided by `mariadb` exists. In addition, we are constrained at Fermilab to avoid distributing servers wherever possible due to the obligation to follow onerous update policies.

Status

[PR7729](#) was submitted 2018-04-11 and became moribund after being nominally approved. It was eventually merged 2019-03-15. A subsequent [issue 11226](#) was addressed and the fix merged upstream as [PR 11237](#).

4.2 ~~Client-only postgresql build~~

As above, server packages are problematic. A build of `postgresql` without the server component is desired.

Status

[PR 7728](#) was submitted 2018-04-11 and became enmired after questions were raised over the appropriateness of virtual packages and separate implementations rather than a (say) `~server` variant. In addition, the PR included a recipe for `py-psycopg2` which was extracted and submitted as a separate [PR 9926](#) by a third party 2018-11-22 and merged 2018-11-24. [PR 7728](#) was closed 2019-03-13 and a revised [PR 10877](#) was submitted with a `client_only` option to change the

way that `postgresql` is built rather than provide a separate client-only product. PR 10877 was merged 2019-03-25.

4.3 Updated root recipe

Our local root recipe had significant changes with respect to the upstream version including more versatility in configuration and use of our local forks of `postgresql-client` and `mariadb-client`. Significant changes to the upstream version were submitted by CERN as PR 8428 2018-06-08 and accepted 2019-02-08 with which our local version is incompatible.

Status

After resolving a problem building more recent versions of ROOT with PR 11129, our local recipe was eventually reconciled with upstream and submitted as PR 11215 with a companion PR 11214 with necessary updates to the recipe for FTGL. PR 11214 was merged 2019-04-23, and PR 11215 was merged 2019-04-24.

4.4 Propagate C++ attributes through dependencies of C++-using packages

Several C++ recipes have a `cxxstd` variant, and depend on other C++ packages. The variant should be propagated in-recipe rather than as part of a `spack install` command, as currently.

Status

A survey needs to be carried out and the appropriate changes to recipes made and propagated upstream.

5 Version History

Version 0.1 (2019-02-25)

First released version.

Version 0.2 (2019-03-19)

- Add table of contents.
- Expand the section on [updates to the version comparison system](#) with examples of [closed and half-open interval version ranges](#), and with a new section on the [ordering issue between versions X and X.Y](#).
- New section, *Compact variant propagation via `depends_on()`*.
- New section, *Improvements to the variants system*.
- Expand section on *Support for external Spack commands* with [contribution from Ben Morgan](#).

- Updates to status for *Client-only mysql-ish packages*, and *Client-only postgresql build*.
- Add section number for *Client-only postgresql build*.
- Add version history.

Version 0.3 (2019-04-30)

- Minor re-word of *Introduction* for clarity.
- Updates to status for:
 - *Support for external Spack commands* (and fix a typo in the title).
 - *The ability to layer Spack installations...*
 - *A persistent relationship between a built package...*
 - *Support for compiler-agnostic packages*.
 - *Base build location on stage...*
 - *New environment manipulation functions...*
 - *Correct ordering between versions X and X.Y* (and expand description).
 - *Client-only mysql-ish packages*.
 - *Client-only postgresql build*.
 - *Updated root recipe*.
- Add status section for *Improvements to the variants system*.
- Consolidate status updates into a single coherent statement of status where appropriate.
- New section, *self.stage.source_path should be usable before it exists*.
- Improve code snippet sections for line overruns and Python-style formatting.

Version 0.4 (2019-05-16)

- New section *Masking of installed packages and dependents*.
- Updates to status for:
 - *New environment manipulation functions...*
 - *No-source packages*.
 - *Configurable equivalency between operating systems*.
- Strikethrough notation is used in item titles to indicate resolution.

Version 0.5 (2019-06-11)

- New section *Install only specified packages from YAML specifications*.
- New section *Writing of the YAML spec file at the start of the build process...*
- Update to status for *self.stage.source_path should be usable before it exists*.