Spack Work Items for HEP

Chris Green greenc@fnal.gov   Lynn Garren garren@fnal.gov
Patrick Gartung gartung@fnal.gov

version 0.3 (2019-04-30)

Contents

1 Introduction

While working on the Spack-based build, development, and release system for HEP, we have come across many instances big and small where Spack core or its recipes require enhancements to allow us to meet our project requirements. In some cases these are showstoppers, such as support for multiple releases and layered (central vs local) installations; in others, they improve an aspect of Spack operation without being essential, such as providing repo-specific package base classes.

This document is an attempt to be a descriptive list of these work items (the most impactful and necessary core items are listed first), along with (where

appropriate) ideas for implementation, statement of importance, and current status.

2 High-Impact Core Enhancements

2.1 Support for external Spack commands

2.1.1 Basic Support for external Spack commands

In order to support the efficient operation of Spackdev, allowing the use of Spack functions and structures internally, `spack dev` should be a spack command. For purposes of ownership and ease of maintenance, this would ideally be an external command, analogous to (*e.g.* `git flow` as an external `git` command).

2.1.2 Enhancement of Spack subcommand facilities

Spack's existing handling of subcommands should be generalized to work for external commands and at any level.

2.1.3 Support for flake8 checks of external Spack commands

A simple way to check external Spack command code for rule-correctness is strongly desired.

2.1.4 Support for external commands and support code in external package repos

The ability to find external commands in external package repos. This is related to a Spack Core Enhancement described elsewhere herein.

Status

PR 8612 from Massimiliano Culpo implemented basic support for external Spack commands. It is now integral to the operation of SpackDev and was merged into the FNAL fork on 2018-09-19. The PR was merged upstream on 2019-03-28 with some changes from the original PR submission as merged into the FNAL fork, which was subsequently reconciled with upstream.

A git branch consisting solely of our proposed enhancement of the subcommand facilities was submitted as PR 11145. Comments were received and incorporated. Subsequently, a competing PR 11209 was submitted by the original author of PR 8612 purporting to address one of the issues addressed by PR 11145. This PR was incorporated into PR 1145 and is awaiting further comments.

While PR 8612 supports testing of external Spack commands, it appears not to support intuitive coding rule enforcement for same. An enhancement request should be discussed with the author of PR 8612.

2.2 The ability to layer Spack installations in a hierarchy

In order to be able to combine locally-built and centrally-installed packages, it becomes necessary to be able to refer to packages in a different Spack installation from the current one.

Status

Originally conceived by Jim Amundson (as "Spack Chains") and entered as PR 5014, and subsequently resubmitted against a more recent version of Spack develompent by Patrick Gartung as PR 8545, this feature has been taken up by Peter Scheibel—a Spack principal— as PR 8772. This PR was merged on 2019-03-27 and a subsequent PR 11152 adding documentation was merged 2019-04-10.

2.3 A persistent relationship between a built package and the options used to build it

In order to support the use of a given Spack installation to support multiple installed software distributions over a significant period of time, it is necessary to be able to identify an installed or cached binary package such that it remains possible to use it even after updates to the recipe(s) have caused a change to the already-built package's calculated hash.

It is the nature of HEP software installations to have a relatively static central installation of software packages and distribution releases against which higher-level packages are developed by users. It must be possible to install centrally newer distribution releases while not compromising the ability of users to build and develop their own software against older ones without unnecessary rebuilds of lower level packages due to hash changes resulting from updated recipes.

Status

The recent Spack Environments functionality may assist with this (as may Spack Chains, described elsewhere in this document), but this remains a thorny problem that will require discussion and exploration before a solution is implemented. The subject has been discussed with Spack principals such as Todd Gamblin and Peter Scheibel and we have been assured that this feature is compatible with the philosophy and direction of Spack, that it has other desirous parties than us, and that it would be reasonable to expect something to look at that should do the job by the end of September 2019.

2.4 Support for compiler-agnostic packages

The HEP community has several applications, distributions of which must be accompanied by substantial data. Examples are particle interaction descriptions and cross-section table. Examples outside our field might include star catalogs

and other astrophysics data. Information of this nature can be quite voluminous (≥ 5~GiB). We would wish to avoid unnecessary duplication for multiple compilers and versions by having just one compiler-agnostic package containing the data.

In addition, pure C and pure interpreted-language packages such as `fftw` or `pyyaml`—or those which combine same—are similarly independent of compiler and version thereof (or perhaps require only a generic compiler runtime such as `libgcc_s.o`). A way to avoid unnecessary duplication of these packages also would be desirable, although the win here would be less in space and more in CPU time for builds.

Status

The subject has been discussed with Spack principals such as Todd Gamblin and Peter Scheibel, including at the telecon 2019-04-11, but no formal enhancement request has been made. Something providing this functionality is expected to be part of the upcoming concretization improvements, but it should be noted that merely allowing the compiler as a "normal" dependency is insufficient to provide the needed flexibility.

2.5 No-source packages

In our existing system, we make significant use of no-source, or "umbrella" packages, whose only function is to facilitate the build or setup of a particular collection of packages.

Status

This feature has been mentioned to Spack principals; no action has been taken.

2.6 Base build location on stage rather than expanded source location

The definition of the CMake build area should be changed from being relative to the staged source directory to being relative to the stage directory.

We have use cases where a recipe needs to have information about the build directory in the environment set by `setup_environment()` (*e.g.* if tests need to have a directory in `PATH`). Specifically, staging happens *after* the call to `setup_environment()`, and if the stage area does not exist, then `spec.stage.source_path` returns, `None`.

Status

PR 8431 was submitted 2018-06-08 and was moribund for several months. This PR was eventually accepted and merged upstream on 2019-03-26.

2.7 `self.stage.source_path` should be usable before it exists.

There are circumstances where it is desirable to *e.g.* set an environment variable in `setup_environment()` based on `self.stage.source_path`. Currently this is not possible as its existence is enforced and `setup_environment()` is executed prior to any staging of the source.

Status

Currently workaround are used when necessary: adding to the default environment for the `make()` object, for example. This enhancement is at the conceptual stage only.

3 Other Spack Core Enhancements

3.1 The ability to inherit from package-repo-specific package base classes

Fermilab has a large number of experiment-related packages using a common CMake-based system called `cetmodules`. Recipes for these packages share a significant amount of boilerplate which could profitably be part of a package base class situated in the same package repo.

Status

This feature has been discussed informally with Spack principals.

3.2 New environment manipulation functions for system path deprioritization and duplicate removal

In order to prevent contamination of `PATH`-like variables with system paths due to `packages.yaml` entries, several new functions provide facilities for `PATH`-like variable manipulation.

Status

PR 8476 was submitted 2018-06-14 and was moribund for a long time. It has been updated against upstream and discussed with Spack principals and further information has been added to the PR. Several changes were necessary to accommodate upstream developments and as of 2019-04-30 it is awaiting final assent and merge.

3.3 Configurable equivalency between operating systems

Due to a change in how Scientific Linux installations identify themselves between SL7.5 and SL7.6, there needs to be a configurable equivalency between compatible

operating systems *e.g.* in `config.yaml`:

```yaml
equivalent-os:
  rhel6:
    - scientific6
    - centos6
  rhel7:
    - scientific7
    - centos7
```

Status

We have a hard-wired equivalency between `rhel7` and `scientific7` in our local code.

A concept for the solution exists (see above). Analysis is needed to determine feasibility and consequences for (*e.g.*) hash values, buildcache packages, *etc.*

3.4 Updates to the version comparison system

3.4.1 Support correct ordering of alpha, beta, pre, rc and patch versions

While it is not usual to need to refer to non-mainstream versions in a recipe, it is certainly not beyond imagining. Several packages have had long release candidate cycles, for instance, and a build of one might be necessary to address an observed issue prior to the final release becoming public. When non-mainstream versions are used in ranges, the version comparison needs to be aware of the relative ordering—release candidates following "pre" versions, but preceding their corresponding released versions, for example.

3.4.2 Support for half-open and closed intervals for version ranges

e.g. describing something which is valid for $7.0 \leq$ `version` < 8.0. For purely numeric versions, the usual workaround of `when=@7.0:7.999` suffices, but for something which is valid also for (*e.g.* release candidates), it becomes more desirable to be able to specify "everything before . . ." Example nomenclature:

```python
# No support before ...
conflicts('cxxstd=17', when='@:<1.27.0')
# Blacklist range.
variant('bad_idea', default=True, when='@0.5.0:<1.0.0')
# Last known good.
conflicts('went_bad_after', when='@1.5.7>:')
# Ignore 2.0.0-alpha, pre, etc.
conflicts('Shaky prerelease', when='@1.8.12>:<@2.0.0'
```

3.4.3 Correct ordering between versions `X` and `X.Y`

There appears to be a problem ordering versions with different levels of speci-
fication. For instance, in `intel-tbb`, the version version of any year's releases
is the unadorned year (`2019`, say), with `2019U1` being the next. These have to
be encoded currently as `2019.0` and `2019.1`, because `2019` does not appear to
compare less-than `2019.1`. Following discussion with Peter Scheibel, it appears
that this is intentional behavior in the numeric case, and additionally applies to
non-numeric versions where one is a front-anchored substring of another (*e.g.*
MVP1 vs MVP1a).

3.4.4 `version()` should recognize filename extensions from `url_from_version()` including `tbz2`

Several of our recipes are of the form where the following `url_from_version()`
is appropriate:

```python
def url_for_version(self, version):
    url = 'http://cdcvs.fnal.gov/cgi-bin/git_archive.cgi/'\
          'cvs/projects/{0}.v{1}.tbz2'
    return url.format(self.name, version.underscored)
```

Unfortunately, for every call to `version()`, it is necessary to specify the extension
explicitly:

```python
version('2.3.0',
        sha256='4b6a29443b631957ca2a7712b5c57762'
        '0c6543e542ee9c77d246cef1e10f7324',
        extension='tbz2')
```

The system should be trained to understand `tbz2` as a valid extension when
returned as a URL by `url_for_version()`.

Status

These core enhancements are at the conceptual stage only. No implementation
has been attempted, but it is likely that they are perfectly feasibly implementable
by a contributor rather than a core expert. This could be entered instead as
an enhancement request, but the likelihood of a core developer taking it on is
somewhat low. The truncated-version-ordering issue must be addressed with
care as there are use cases for both behaviors, and they should be distinguishable
and selectable as appropriate.

3.5 Compact variant propagation via `depends_on()`

There is a need for propagation of variants to dependencies generally, rather
than specifically. For instance, in order to pass the value of the `cxxstd` variant
down a dependency tree, currently one must do:

```
depends_on('boost cxxstd=default', when='cxxstd=default')
depends_on('boost cxxstd=98', when='cxxstd=98')
...
```

One could imagine instead a more compact syntax, something like:

```
depends_on('boost', propagate_variants=['cxxstd', 'python'])
```

If `cxxstd` is not specified, then nothing will be propagated (if the variant exists in the dependency it will be defaulted). Otherwise, the value of `cxxstd` in the current recipe will be used.

As a further extension, using `recursive_variants` instead of `propagate_variants` would cause the specified variants to be propagated to lower-level dependencies *even if intermediate dependencies did not have the variant(s) in question defined.*

Status

These enhancements are at the conceptual stage only.

3.6   Improvements to the variants system

Several packages have extensive feature options, which evolve over time. Features come and go, and depend upon each other. An expanded variant specification syntax is desired to be able to specify these variants and the relationships between them. The archetype for this use case would be the `boost` recipe. Some ideas:

3.6.1   `variant()` grows a `when` clause

The `when` clause shall specify when the variant is defined. It shall be valid therein to refer to versions and to any other variants that have already been defined.

3.6.2   New function `variant.depends_on()`

This function, having a `when` clause as defined for `variant()`, above, shall specify upon what other variant(s) `self` shall depend when the `when` clause is valid. It shall return the variant object allowing for chaining, *e.g.*

```
variant('context', when='@1.61.00:').\
  depends_on('thread', 'chrono', 'system', 'date_time')
variant('fiber', when='@1.62.00:').depends_on('context').\
  depends_on('thread', 'chrono', 'system', 'date_time',
             when='@:1.68.99')
variant('signals', when='@1.29.00:1.68.99')
```

Status

These enhancements are at the conceptual stage only.

4 Recipe Enhancements

4.1 Client-only `mysql-ish` packages

The existing `mariadb` package is GPL and therefore not usable by HEP community in a linking context. An LGPL client-only package provided by `mariadb` exists. In addition, we are constrained at Fermilab to avoid distributing servers wherever possible due to the obligation to follow onorous update policies.

Status

PR7729 was submitted 2018-04-11 and became moribund after being nominally approved. It was eventually merged 2019-03-15. A subsequent issue 11226 was addressed and the fix merged upstream as PR 11237.

4.2 Client-only `postgresql` build

As above, server packages are problematic. A build of `postgresql` without the server component is desired.

Status

PR 7728 was submitted 2018-04-11 and became enmired after questions were raised over the appropriateness of virtual packages and separate implementations rather than a (say) `~server` variant. In addition, the PR included a recipe for `py-psycopg2` which was extracted and submitted as a separate PR 9926 by a third party 2018-11-22 and merged 2018-11-24. PR 7728 was closed 2019-03-13 and a revised PR 10877 was submitted with a `client_only` option to change the way that `postgresql` is built rather than provide a separate client-only product. PR 10877 was merged 2019-03-25.

4.3 Updated `root` recipe

Our local `root` recipe had significant changes with respect to the upstream version including more versatility in configuration and use of our local forks of `postgresql-client` and `mariadb-client`. Significant changes to the upstream version were submitted by CERN as PR 8428 2018-06-08 and accepted 2019-02-08 with which our local version is incompatible.

Status

After resolving a problem building more recent versions of ROOT with PR 11129, our local recipe was eventually reconciled with upstream and submitted as PR 11215 with a companion PR 11214 with necessary updates to the recipe for FTGL. PR 11214 was merged 2019-04-23, and PR 11215 was merged 2019-04-24.

4.4 Propagate C++ attributes through dependencies of C++-using packages

Several C++ recipes have a `cxxstd` variant, and depend on other C++ packages. The variant should be propagated in-recipe rather than as part of a `spack install` command, as currently.

Status

A survey needs to be carried out and the appropriate changes to recipes made and propagated upstream.

5 Version History

Version 0.1 (2019-02-25)

First released version.

Version 0.2 (2019-03-19)

- Add table of contents.
- Expand the section on updates to the version comparison system with examples of closed and half-open interval version ranges, and with a new section on the ordering issue between versions X and X.Y.
- New section, *Compact variant propagation via `depends_on()`*.
- New section, *Improvements to the variants system*.
- Expand section on *Support for external Spack commands* with contribution from Ben Morgan.
- Updates to status for *Client-only `mysql`-ish packages*, and *Client-only `postgresql` build*.
- Add section number for *Client-only `postgresql` build*.
- Add version history.

Version 0.3 (2019-04-30)

- Minor re-word of *Introduction* for clarity.
- Updates to status for:
    - *Support for external Spack commands* (and fix a typo in the title).
    - *The ability to layer Spack installations. . . .*
    - *A persistent relationship between a built package. . . .*
    - *Support for compiler-agnostic packages*.
    - *Base build location on stage. . . .*
    - *New environment manipulation functions. . . .*
    - *Correct ordering between versions X and X.Y* (and expand description).
    - *Client-only `mysql`-ish packages*.
    - *Client-only `postgresql` build*.

- *Updated `root` recipe.*
- Add status section for *Improvements to the variants system*.
- Consolidate status updates into a single coherent statement of status where appropriate.
- New section, *`self.stage.source_path` should be usable before it exists*.
- Improve code snippet sections for line overruns and Python-style formatting.